

The Unified Framework STOCK RESEARCH Suite v2.1

VERSION 2.1 | MULTIBAGGER CAPABILITY RELEASE | SEPTEMBER 2026

Supersedes Unified Framework v2.0 in full.

Component versions under this release: Stage 0 v1.3 · Stage 1 v3.5 · Stage 2 v7.0 · MQF v1.1, plus a new Part 7 — The Multibagger Engine integrated across Stage 2 (new Sections 5b-T, 5e, 5f, 6A, 6B) and orchestrated by a consolidated candidate gate (N1, N7, N8).

- >> Diagnosed and corrected: durability now gates ahead of price, not behind it
- >> Adds the return arithmetic the suite never computed — a mandatory Multibagger Return Bridge
- >> Adds the compounding engine test: reinvestment rate $\times$ incremental ROIC, as a binding gate
- >> Moat testing is now directional (BUILDING / HOLDING / ERODING), not a static possession
- >> Rewrites the exit and position-cap rules that mathematically forbade a 10-bagger
- >> Pays its M2 complexity budget by retiring three rules; every new rule enters at PILOT
- >> All Part 0 meta-layer rules, all kill-gates, and all MQF/governance caps carry forward unchanged

NOT investment advice — methodology reference for personal research only — consult a SEBI-registered adviser

Severity Colour Legend (applies to every report produced under this suite)

Colour	Band	Meaning in a report	Consequence
GREEN	PASS / STRENGTH	Evidenced positive: gate cleared, engine test passed, moat trajectory BUILDING, clean governance finding	None — proceed
TEAL	INFORMATION	Neutral structural fact, methodology note, or figure with no directional read	None — context only
ORANGE	MEDIUM CONCERN	Irani AMBER, prudential cap, borderline score, PILOT rule invoked, engine test partially met	Cap, monitor, or one-notch conviction reduction
RED	SEVERE / DISQUALIFYING	Irani RED, kill-gate fired, G3-A, value-trap disqualifier, MQF AVOID, candidate-gate failure	Discard or hard cap — never outvoted by score
NAVY	STRUCTURAL RULE	A binding rule, cap or definition quoted from the framework itself	Not relaxable

Every finding in every report carries exactly one band. A finding with no band is an incomplete finding. Colour is a rule, not decoration.

What changed in v2.1, in one paragraph

v2.0 was a rigorous trap-avoidance engine wearing a multibagger label: Stage 1 ("is it cheap?") gated Stage 2 ("is it durable?"), so durability could only confirm a verdict that cheapness already permitted; ROCE was scored as a level with no reinvestment-rate test, so a high-return, low-payout compounder scored no better than a high-payout mature cash cow; the moat test asked only whether an advantage existed, never whether it was widening; and Part 5's exit rules trimmed a position at base-case fair value and forced a review at 10% of book — exactly where a multibagger's return is made. v2.1 corrects all four. It does not loosen a single governance rule, kill-gate, or MQF cap.

Architecture map

Layer	Question it answers	Status in v2.1
Part 0 — Meta-Layer (M1–M9)	Is the framework itself still trustworthy?	Unchanged from v2.0 — binding on every component below
Part 1 — Stage 0 (v1.3)	Is a corporate-action turnaround being staged, and when can evidence exist?	Amended — S1 tracking carve-out, new S12
Part 2 — Stage 1 (v3.5)	Is it cheap, is it a trap, and at what price?	Amended — trajectory-conditioned MATM, Engine Compounder MoS tier
Part 3 — Stage 2 (v7.0)	Is it durable, of what tier, how large, when to sell?	Amended — Section 1A scoped; five new sections (5b-T, 5e, 5f, 6A, 6B)
Part 4 — MQF (v1.1)	Can the promoter be trusted with the capital?	Unchanged from v2.0 — can still veto Stage 2 sizing outright
Part 5 — Portfolio & Execution	How does this position fit the book, and when does it leave?	Amended — P-1, P-5, P-6, P-7 rewritten to stop capping winners
Part 7 — Multibagger Engine (NEW)	Does this business have a mechanism that can multiply capital, and is that mechanism available at this price?	New — orchestrates the five new Stage 2 sections into one gate
Part 6 — Output, Checklist, Integration	How is all of the above assembled and delivered?	Amended — Engine×Price matrix replaces the Cheap×Durable 2×2

Part 0 — The Meta-Layer (M1-M9)

Unchanged from v2.0 in every particular. Reproduced here in full because this document supersedes v2.0 in full — nothing is left to a superseded PDF. Binding on all components below; cannot be relaxed by any component rule.

M1 Base-Rate Discipline	NAVY
No rule, trigger, band or multiplier may carry status VALIDATED without a stated denominator: N names screened, N that fired the rule, N that subsequently worked, N that failed. Status ladder: PILOT (no denominator) → PROVISIONAL (denominator stated, n<3 successes or unstable hit-rate) → VALIDATED (denominator stated, 3+ independent successes, failures disclosed). Every rule new to v2.1 (Part 7, and the new Stage 2 Sections 5b-T/5e/5f/6A/6B) enters at PILOT. Why / how it binds: Prevents a validation claim from resting on hindsight-selected winners. It is the single rule that keeps this whole suite honest about what it has actually earned.	
M2 Complexity Budget	NAVY
The suite is capped at its current rule count. A new rule may be added only if (a) an existing rule is retired, or (b) the new rule arrives with an M1-compliant denominator. Annual audit deletes any rule that has not changed a verdict in 24 months. Why / how it binds: v2.1 added eight new rules (Section 5b-T, 5e, 5f, 6A, 6B, plus N1/N7/N8) with no denominators available. It pays for them by retiring the Section 1A blanket full-coverage rule, the Rule #26 four-band VIX table, and the Layer A dividend-yield metric for engine-assessed names — see the M2 retirement ledger at the end of this document for the full accounting.	
M3 Edge Hypothesis (must be stated per position)	NAVY
Claimed edges, in priority order: (1) time-horizon arbitrage — holding through drawdowns that force institutional selling; (2) coverage neglect — small/microcaps below institutional research thresholds; (3) behavioural discipline — pre-committed exits and pre-mortems reducing unforced errors. Explicitly NOT claimed: superior data, modelling, speed, or information advantage. Why / how it binds: Every BUY must name which edge it relies on. A name relying on none is capped at WATCH — you have no reason to expect to be right.	
M4 Framework Kill-Switch (two-sided)	NAVY
Downside: at the 24-month check, if BUY-verdict names underperform the relevant index on a cost- and tax-adjusted basis across a majority of closed positions, MANDATORY REVIEW fires — new positions paused. Upside/inaction: rejection rate above 95%, or zero BUY verdicts across four consecutive quarters, fires the same review. Why / how it binds: Answers both failure modes — a process that loses money, and a process tightened until it can never act. The second is the more likely outcome of this suite's amendment history, v2.1 included, which is precisely why this rule stays two-sided.	
M5 Adversarial Protocol (replaces solo self-assessment)	NAVY
Before any first purchase: (a) steelman the short — the strongest bear case, in the bear's own framing; (b) disconfirming-evidence search, logged, including one hostile source; (c) 48-hour cooling period before first purchase, no new buying arguments admitted; (d) the Section 9A pre-mortem, written before the verdict. Why / how it binds: A solo pre-mortem written after the work is confirmatory by default. The steelman and the logged disconfirming search are what can actually change a mind. This protocol also protects the Part 7 N1 return bridge from becoming arithmetic-shaped advocacy — see Part 7.	
M6 Source-Tier Hierarchy	NAVY
Tier 1 audited statements, regulatory orders, court judgments. Tier 2 exchange filings, shareholding patterns, CARO/MR-3. Tier 3 aggregators (Screener, Trendlyne, Tickertape) — must be cross-checked against Tier 1/2 for any gate-driving figure. Tier 4 management narrative (concalls, investor decks, MD&A). Tier 5 media, brokerage notes, short-seller reports — leads only. Why / how it binds: No gate clears on Tier 3 alone. No Section 5d test rates GREEN on Tier 4 alone. Every figure in a report carries its tier — including every figure in the new Part 7 gates.	

M7 Benchmark Hurdle (opportunity cost)	NAVY
Every BUY/ACCUMULATE must state: the relevant index benchmark (Nifty 50 / Midcap 150 / Smallcap 250 by cap band), the required annualised return over the stated horizon, and the specific driver expected to deliver it. Suggested floor: index return plus a risk premium sized to the MoS tier.	
Why / how it binds: A name can clear every gate and still be a worse use of capital than an index fund at near-zero effort. This is the benchmark the Part 7 N1 return bridge must clear net of E-3 costs.	
M8 Funnel Discipline	NAVY
Record the candidate source for every name entering the funnel: systematic screen / Stage 0 trigger / news or media / peer mention / existing holding review / other. Annual audit computes hit-rate and verdict distribution by source. Default minimum 50% of the annual funnel from systematic screens.	
Why / how it binds: Coverage is the binding constraint — a full report is hours of work — so which names enter the funnel matters more than how they are scored. Part 7's N8 Multibagger Funnel Screen is the systematic screen this rule requires but never itself specified.	
M9 Two-Pass Sequencing	NAVY
Stage 1 Pass 1 runs with MATM fixed at 1.00x and produces a provisional Layer A and composite. Stage 2 then runs in full. Once Section 5b classifies the moat (and, under v2.1, Section 5b-T classifies its trajectory), Stage 1 Layer A is re-scored once with the correct MATM anchor. Both composites are logged; Pass-2 governs.	
Why / how it binds: Resolves the circular dependency between Stage 1's MATM and Stage 2's moat classification, and — as of v2.1 — carries the moat-trajectory result forward into the MATM conditioning rule at the same handoff point, at no extra sequencing cost.	

Part 1 — Stage 0: Special Situations & Turnaround, v1.3

Question: is a corporate-action turnaround being STAGED, and when can the evidence exist? Stage 0 never produces a buy verdict, a position size, or a gate exemption — the single narrow exception is the v1.3 tracking-position carve-out at S1, capped at 0.5% and hedged by an automatic sunset.

ID	Trigger	Fire condition	Class	Status (M1)
T1	Scheme of arrangement	NCLT scheme approved or merger/demerger effective date announced	A	PILOT
T2	Name / identity change	Company name change, especially shell-to-brand	A	PILOT
T3	Promoter stake rising	Promoter holding up ≥ 5 pp across ≤ 2 quarters	A	PILOT
T4	Preferential issue to promoters	Warrants/equity subscribed by promoter group	A	PILOT — co-validated, shares only case with T5
T5	Founder / credible promoter return	Founder returns to board/MD, or credible new promoter	A	PILOT — co-validated, shares only case with T4
T6	Demerger / subsidiary event	Record date set; subsidiary listed/delisted; carve-in/out	A	PILOT
T7	First profit after loss streak	First profitable quarter after ≥ 3 consecutive loss years	B	PROVISIONAL — 3 successes, denominator not yet counted
T8	Restatement-driven revenue jump	Revenue $>100\%$ YoY with restatement/merger-accounting note	B	PILOT
T9	Deleveraging complete	Net debt $\rightarrow$ net cash within 4 quarters, or $>50\%$ debt reduction	B	PROVISIONAL — 2 successes only
T10	Re-rating in a small cap	Price $>50\%$ above 52-week low while mcap $< ₹10,000$ Cr	C	PILOT
T11	First institutional entry	Bulk/block deal or new FII/DII holder, no prior history	C	PILOT

Class A = corporate-event, Class B = financial-inflection, Class C = microstructure. At least one clustered trigger must be Class A or B. Two Class C triggers never qualify — that is momentum-chasing, not situation detection.

Clustering, exclusions, re-screen condition

Triggers in rolling 6 months	Result	Action
1	NOISE	No action, log nothing
2, with ≥ 1 Class A/B	STAGE 0 CANDIDATE	Open Turnaround Journal entry; dated re-screen condition; log cluster validation composition (S11)
3+, with ≥ 1 Class A	STAGE 0 PRIORITY	As above, plus pull scheme docs/AR immediately; pre-draft the Stage 1 data file
Spread >6 months apart	RESET	Window is rolling; stale triggers age out
Exclusion	Test	Band
X1 — Fraud-state promoter	Siphoning, funds to promoter entities, CBI/ED arrest — at ANY group entity	RED
X2 — Active PFUTP	Active SEBI PFUTP on sitting MD/CMD/CFO/WTd, not stayed or disposed	RED
X3 — SME / illiquid shell	30-day ADTV persistently below ₹25 lakh AND mcap below ₹250 Cr	RED
X4 — Serial restructurer	Third or later scheme/name change/reorganisation at the same entity within 5 years	RED

Re-screen condition, required at entry, no exceptions: earnings evidence (2 consecutive clean standalone profitable quarters post-event, 3 if merger accounting); governance evidence (all open prudential flags resolved per Stage 1 Layer D); structural evidence (corporate action complete); liquidity evidence (30-day ADTV clearing the Part 5 E-2 sized-position floor); a calendared date. Anti-FOMO (S5): a price rally is never a re-screen trigger — evidence dates only.

Rules S1-S12

#	Rule
S1 (v1.3 amended)	Stage 0 never produces a buy verdict, position size, or attractive-below price, with one narrow exception: a 0.5% tracking position is permitted at a qualifying cluster if ≥ 1 Class A trigger has fired, X1-X4 are all clear, G3 state is D, and G0/E-2 liquidity floors are met at that size. The dated re-screen condition is still logged unchanged and full sizing still waits for the evidence date. Never a gate exemption, a verdict, or grounds to add.
S2	2+ triggers in a rolling 6-month window, at least one Class A/B. Two Class C never qualify.
S3	X1 and X2 exclude absolutely — do not watch, do not journal beyond the exclusion note.
S4	Every entry gets a dated re-screen condition AT ENTRY. No condition, no entry.
S5	Price action is never a re-screen trigger. Evidence dates only.
S6	Consolidated-only profit never satisfies the earnings evidence requirement.
S7	Never re-screen mid-corporate-event.
S8	Maximum 3 re-screen cycles per entry, then archive with reason. The archive retires the ENTRY, not the NAME — a completed corporate action re-admits the name as a fresh Stage 1 candidate.
S9	Stage 0 membership confers NO exemption from any Stage 1/2 gate — front-door re-entry always.
S10	Trigger status follows PILOT/PROVISIONAL/VALIDATED (M1) and requires a denominator for any validation claim. No trigger in the suite currently holds VALIDATED status.
S11	Every journal entry records the validation composition of its cluster, e.g. 'T4 [PILOT] + T3 [PILOT]', for later false-negative review weighted by evidence strength.
S12 (v1.3 NEW)	Every S1 tracking position is tagged PILOT under M1 with its own denominator track: N opened, N reaching the evidence date intact, N exited at a loss first. Reviewed at the annual audit; a negative 24-month record reverts S1 to its pre-v1.3 absolute text automatically.

Flagged for attention: S1/S12 is the single most dangerous rule change in this release — it puts a hole in a rule that worked precisely because it was absolute. The 0.5% cap, the Class-A requirement and the self-reverting sunset exist to bound the damage; if any one amendment here should be rejected on review, it is this one.

Part 2 — Stage 1: Undervalued Stock Identification, v3.5

Question: is it genuinely cheap, is it a trap, and at what price? Runs Pass 1 at MATM 1.00x, then re-scores once after Stage 2 classifies the moat and — new in v2.1 — its trajectory (M9).

Gate	Test	Consequence	Band
G0 Liquidity	30-day ADTV below ₹1 Cr; also tests the Part 5 E-2 sized-position floor	DISCARD	RED
G1 Leverage	D/E above 2.0x AND interest coverage below 1.5x simultaneously	DISCARD	RED
G2 Earnings	Negative PAT for 3 consecutive years (excl. regulated utilities, BFSI, turnarounds)	DISCARD unless turnaround	RED
G3 Promoter integrity	Four-state A/B/C/D taxonomy	State-dependent cap	RED/ORANGE
G4 Objectivity	Enforced by the M5 adversarial protocol, not self-assessment	Protocol incomplete = cannot lock	ORANGE
G5 Valuation sanity	P/E >3x sector median → REVIEW. Eliminates only if all four G5 conditions hold	INVESTIGATE	ORANGE
G3 state	Definition	Cap	Band
G3-A	Siphoning; funds to promoter entities; CBI/ED arrest	0% — full discard	RED
G3-B	Pledge 10–25%; promoter holding <15%; settled enforcement	Max 1.5%	ORANGE
G3-C	Active SEBI PFUTP on sitting MD/CMD/CFO/WTD	Max 1%	RED
G3-D	Pledge <10%; no investigation; holding >25%; clean	Standard	GREEN

Irani governance flags

Flag	Trigger	Position impact	Band
Accounting Gap	PAT>OCF 3+yrs; receivables outrunning revenue; exceptionals >50% PAT	Layer C EQ zeroed; value-trap flag	RED
SEBI PFUTP	Active, not stayed/disposed	G3-C; position max 1%	RED
Leverage Trap	D/E>1.5x rising AND interest coverage<2.0x	G1 logic; value-trap flag	RED
Founder Large OFS IPO	OFS>60% AND fresh issue<40%; founder seller	Layer D –2; cap 3%	ORANGE
Value Extraction	Foreign PE promoter + payout>100% for 2+yrs + rev CAGR<10%	Layer D capped 8/15	ORANGE
Di-worse-ification	Frequent unrelated acquisitions; capital destroyed in non-core	Layer C –3; Layer D –2	ORANGE
Commodity Dependence	Revenue 1:1 with commodity; COGS>70%, GM<15%	Passthrough discount	ORANGE
ROCE Decay	ROCE down >3pts/2yrs, no capex explanation	Mandatory review	ORANGE
Correlation Concentration	Single macro factor >30% of portfolio	Enforced at Part 5 P-3	ORANGE
Auditor/CFO Resignation	Outside routine rotation, or CFO/COO departs w/o succession, within 18mo	Layer D –2; explanation required before BUY	ORANGE — prudential

Section 4A — Rule #25

Step 1 — basis. FORWARD EPS at 75% utilisation only if BOTH hold, evidenced separately: (a) the G5 trough/capex/turnaround definition; AND (b) a completed, dated Section 4B bridge. Absent (b), TTM EPS applies even where (a) holds. Steady-state names are always TTM.

Outcome	Test	Consequence
PASS	Earnings yield (correct basis) $\geq$ live 10-yr G-Sec, sourced and dated	Full scoring; M7 hurdle still applies separately
REVIEW BAND	Fails on yield alone, but (earnings yield + sustainable dividend yield) $\geq$ G-Sec AND fwd P/E $< 1.2\times$ sector median	Scoring proceeds, capped at ACCUMULATE/STEADY COMPOUNDER — definitional, not relaxable
FAIL	Earnings yield fails and the total-yield test also fails	WATCH only

Mandatory caveat, printed verbatim on every REVIEW-BAND invocation: "[PILOT — n=1] The Total-Yield Tolerance Band rests on a single validation case (ITC), used as both calibration and live invocation. UNVALIDATED, not confirmed. Conviction reduced one notch." Nominated for deletion at the next annual audit unless a third, genuinely independent case is logged first — see the M2 retirement ledger.

Layer A — relative valuation (30 pts)

Metric	How tested	Attractive when
P/E — TTM / 1yr fwd / 2yr fwd	Forward EPS anchored to capacity/guidance, never hope	Below own 5-yr and sector median on TTM, or compelling forward
PEG	3-yr forward EPS CAGR, not historical	Below 1 attractive; ~ 1 fair; above 1.5 rich
P/B (sector-adjusted)	Sector thresholds table	Below sector deep-value band
EV/EBITDA — TTM + fwd	Forward at 75% utilisation; mandatory for capex-phase	Below peers, debt-neutral
FCF yield	Operating FCF; capex-phase years documented and excluded	Above 5–7%
Earnings yield vs G-Sec (#25)	Section 4A basis logic + tolerance band	Comfortably above G-Sec on the correct basis
MATM (bounded, v2.1 trajectory-conditioned)	Exactly one of 1.00x(none)/1.10x(NARROW)/1.25x(MODERATE)/1.50x(WIDE), read off Section 5b. v2.1: the 1.10/1.25/1.50 anchors apply only where Section 5b-T rates trajectory BUILDING or HOLDING. ERODING caps MATM at 1.00x regardless of moat class. Pass 1 always uses 1.00x.	Current multiple below the conditioned MATM fair multiple
Dividend yield — retired for Part 7 engine-assessed names (v2.1, D-2/N2-4)	Payout sustainability; feeds the tolerance band. Kept in full for the non-engine value path.	Above own history and G-Sec — value path only
M7 Benchmark hurdle	Required annualised return vs the relevant index, driver named	Expected return clears index plus tier-appropriate risk premium

Section 4C — intrinsic value, EFV, bear-case gate

Two-phase DCF mandatory for ramp/capex-phase businesses (Phase 1 ramp CAGR, Phase 2 steady-state — never a single blended CAGR). Probabilities default 25/50/25, pre-registered with a timestamp before bear/base/bull values are computed; deviation requires a specific dated external fact. $EFV = (P_{bear} \times Bear) + (P_{base} \times Base) + (P_{bull} \times Bull)$, reported alongside the base case, never replacing it. Quantitative bear-case gate: probability-weighted downside must not exceed the Required MoS for the tier. UNDERVALUED additionally requires bear-case forward P/E $\leq 2\times$ sector median.

Business quality tier	ROCE / character	Required MoS	Band
Exceptional compounder	Above 25%, wide moat, clean governance	15–20%	GREEN
Good quality	15–25%, moderate moat	25–30%	GREEN
Average / improving	10–15%, improving trend	35–40%	TEAL
Cyclical / commodity	Highly variable, commodity-linked	40–50%	ORANGE
Turnaround / distressed	Negative to low, recovery thesis	50%+ — size very small	ORANGE
Capex-phase ramp	Below target, ROCE recovery documented	25–35% on forward intrinsic	TEAL

Business quality tier	ROCE / character	Required MoS	Band
Engine compounder (v2.1 NEW)	ROIIC>20%, RR>50%, trajectory BUILDING, all Part 7 gates passed, MQF HIGH-TRUST (85+) — all five required, none waivable	10-15% on the N1 bridge, taken in duration: bridge must still clear M7 if high-return duration is cut to 60% of assumed period	GREEN

Abuse guard on the Engine compounder row: gated on MQF HIGH-TRUST — the strictest governance band in the suite — precisely because a compounder-grade MoS on an ungoverned business is how value frameworks blow up. Any MQF band below 85 routes the name back to the tiers above.

Composite, verdict bands, sizing

Layers: A relative valuation /30 · B intrinsic & MoS /30 · C business quality /25 · D governance & safety /15.
Round to nearest 5; within ± 3 of a boundary is BORDERLINE, resolved by a Layer C+D tie-break, never by the decimal.

Composite	Verdict	Band
80-100	STRONGLY UNDERVALUED — high-conviction buy at the appropriate MoS	GREEN
65-79	UNDERVALUED — accumulate on dips; confirm catalyst	GREEN
65-79 + fwd PE compelling	STEADY COMPOUNDER — route to Stage 2, accumulate slowly	GREEN
50-64	FAIRLY VALUED — watchlist; re-screen quarterly	ORANGE
35-49	SLIGHTLY OVERVALUED — no action / trim	ORANGE
Below 35	OVERVALUED — avoid / exit on strength	RED
Any disqualifier	VALUE TRAP FLAG — reject regardless of score	RED

Condition	Tier	Size
Score 80-100, Layer D clean, 2+ catalysts	HIGH CONVICTION BUY	5-7%
Score 65-79, Layer D clean, 1+ catalyst	BUY / ACCUMULATE	3-5%
Score 65-79 + fwd PE compelling, 1+ catalyst	STEADY COMPOUNDER	2-4%
Rule #25 REVIEW-BAND pass, any composite	Capped — not relaxable	Max applicable tier; never 5-7%
Score 65-79, 0 catalysts	WATCH — no catalyst	0-1%
M3 edge not nameable / M7 hurdle not cleared	Capped at WATCH	0-1%
G3-C / G3-B / G3-A governance caps	Binding governance caps	1% / 1.5% / 0%
4+ conditional flags	Exception-stack cap (prudential)	Max 1%

Cap Reconciliation: GOVERNANCE caps (G3-A/B/C, PFUTP, siphoning, Founder-OFS, Value Extraction, MQF cap) are binding, never relaxed. PRUDENTIAL caps (exception stack, catalyst tier, volatility) manage timing/uncertainty and may not alone force an intact high-conviction thesis below its documented floor. DEFINITIONAL caps (Rule #25 REVIEW band, M3, M7) are part of the rule that creates them and never relaxable.

Part 3 — Stage 2: The Multibagger Deep-Dive, v7.0

Question: is this a durable compounder, of what tier, how large, and when to sell? v2.1 adds five new sections that give this stage the ability to actually recognise a compounding engine, not merely a moat.

Section 1A — Full-Coverage Rule (v2.1 amended, closes D-8)

v2.0 text (superseded): full Sections 5–9A and Modules A/B/C ran on every name Stage 1 screened, including every AVOID and WATCH. This directly conflicted with M8, which identifies analyst hours as the binding constraint and the funnel as the dominant driver of outcomes — Section 1A was spending that constraint on rejected names.

v2.1 text (governs): full Stage 2 runs on every name Stage 1 scores FAIRLY VALUED or better, and on every name Stage 1 rejects on price alone with Layers C+D otherwise clean (the multibagger-relevant rejection class, preserved deliberately). Every other rejection — governance-driven, quality-driven, or trap-flagged — receives a 10-line qualitative note plus the mandatory M5 pre-mortem, no Module numbers. Sole carve-out unchanged: G3-A discards get the qualitative note only, no scoring at all, since the financials are themselves suspect. Sizing independence is unchanged — a high Module score on a Stage-1-rejected name is calibration only, never a route back to BUY.

Section 5 — moat analysis (Module A foundation)

Moat source	Evidence required	Rating logic
1. Pricing power	5-yr gross margin trend; price-hike history vs volume	GREEN: GM stable/rising through downturns. RED: no pricing power, GM<20%
2. Switching costs	Integration depth; churn; contract length	GREEN: multi-year plus integration. RED: commodity, easy switch
3. Network effects	User CAGR vs revenue CAGR; platform metrics	GREEN: demonstrated with data. RED: none
4. Cost advantage	EBITDA margin vs peers; named process/technology	GREEN: >300bps AND innovation-driven. Scale-only capped AMBER
5. Intangible assets	Documented price premium; patent/licence uniqueness	GREEN: documented premium or unique licence
6. Efficient scale	Market size; entrant ROCE vs incumbent	GREEN: duopoly/oligopoly, entrant ROCE below cost of capital

Class	Sources	GREEN ROCE band	Module A range	MATM anchor
NONE	0–1	Below 12%	0–20	1.00x
NARROW	1–2	12–18%	20–45	1.10x
MODERATE	2–3	15–25%	45–65	1.25x
WIDE	3+	Above 25%	65–100	1.50x

A moat premium is earned by evidence, never assertion. Under M6 the evidence must be Tier 1–3; Tier 4 narrative alone cannot make a source GREEN.

Section 5b-T Moat Trajectory (v2.1 NEW — closes diagnosis D-5)

NAVY

Every Section 5 moat source rated GREEN or AMBER is additionally classified BUILDING / HOLDING / ERODING over 3–5 years, on evidence: market-share trend; gross-margin trend versus the sector's; price realisation versus input costs; customer count and retention; the company's share of the industry's incremental capacity additions; brand/R&D/distribution spend and its measurable output.

Rule (binding, feeds Stage 1 Layer A MATM): the v2.0 MATM anchors 1.10/1.25/1.50 apply only where trajectory is BUILDING or HOLDING. ERODING caps MATM at 1.00x regardless of moat class. Multibagger Candidate tier (Part 7) requires BUILDING on at least one source, with Tier 1–3 evidence. HOLDING throughout caps at Steady Compounder. Re-rated annually alongside Rule #27; a BUILDING→HOLDING transition is a logged conviction event, not a silent one.

Why / how it binds: v2.0 grades the moat a company has; the return comes from the moat it is acquiring. A priced-in moat is what the market is best at valuing and therefore the worst place to look for a mispricing. A widening moat is what it is worst at valuing, because widening requires believing a trend will persist — which is exactly the kind of belief a coverage-neglected smallcap can hold well before the market catches up.

Section 5c — Rule #27 moat decay scorecard (mandatory annual)

Erosion signal	Trigger	Impact	Band
Pricing power loss	GM down >200bps YoY for 2 years	Moat –2	ORANGE
Market share erosion	Revenue CAGR below sector for 3 years	Moat –2	ORANGE
Technology disruption	>20% of TAM threatened by new technology	Moat –3	RED
Customer concentration	Top-3 customers >50% AND one exits	Moat –2	ORANGE
Regulatory moat removed	Licence/patent/regulatory barrier lapsed	Moat –4; full re-score	RED
ROCE decay	Down >3pts/2yrs, no capex explanation	Links to Section 8	ORANGE

Any RED = mandatory thesis review. Two or more AMBER = reduce conviction tier one step.

Section 5d — Innovation and Learning Culture (mandatory)

Test	Question	Evidence floor (M6)
I-1 Value chain innovation	Structural manufacturing/logistics/procurement improvement in 3–5 yrs?	GREEN requires Tier 1–3 corroboration; concall claims alone cap AMBER
I-2 Digital/technology adoption	Creating differentiation, or merely meeting compliance?	Same floor — measurable output, not a stated initiative
I-3 Regulatory adaptability	Speed/effectiveness in adapting to regulatory change historically	Tier 1–2 available by nature — no excuse for Tier 4 here
I-4 Workforce & succession depth	Institutional knowledge beyond founder/current MD? Real second line?	Named executives with tenure/disclosed responsibility, not org-chart assertion

Scoring: 3+ GREEN = STRONG. 2 GREEN+1–2 AMBER = ADEQUATE. 2+ RED = WEAK (Module A –3, Module B runway capped). v6.5 cross-check: a STRONG/ADEQUATE rating against 2+ Rule #27 AMBER signals, or against declining forward ROCE with no capex explanation, is downgraded one tier and the conflict named in the Section 9A pre-mortem.

Section 5e The Compounding Engine — reinvestment rate × incremental ROIC (v2.1 NEW, binding gate)

NAVY

Compute and report three figures: $ROIIC = \Delta NOPAT (3\text{-yr}) \div \Delta \text{invested capital} (3\text{-yr})$, where invested capital = net fixed assets + CWIP + net working capital; $\text{Reinvestment rate} (RR) = (\text{capex} + \Delta WC - \text{depreciation}) \div NOPAT$; $\text{Engine growth } g = ROIIC \times RR$.

Gate 1 (tier gate). Multibagger Candidate tier requires $g \geq$ the required net-profit CAGR from the Part 7 N1 return bridge, computed on ROIIC, not headline ROCE. Gate 2 (floor). ROIIC below cost of capital for two consecutive years is RED — value is being destroyed at the margin while headline ROCE lags on legacy assets. Gate 3 (payout conflict). 3-yr average payout above 40% forecloses Multibagger Candidate tier outright; route to Steady Compounder. High payout and high compounding are arithmetically incompatible. Gate 4. Layer A's dividend-yield metric is retired for names assessed under this section — see the M2 retirement ledger.

Why / how it binds: This is the single largest omission the suite carried before v2.1, and the reason its output skewed toward mature, cash-returning, optically cheap businesses. ROCE tells you what the past capital earned; ROIIC tells you what the next rupee will earn; the reinvestment rate tells you how many rupees there will be. A multibagger is a high ROIIC compounded at a high reinvestment rate for longer than the market assumed — nothing more exotic than that.

Section 5f Promoter Alignment (v2.1 NEW — forward-looking, distinct from MQF)**NAVY**

MQF answers whether the promoter is honest. This section answers whether the promoter is building. Report: promoter holding level and 3-year direction; open-market promoter purchases (Reg 29/31, M6 Tier 2); pledge trend; the 5-year record of capital deployed and the ROIIC it earned (links to Section 5e); ESOP structure and whether vesting is performance-linked; related-party trend (inherits MQF Section 7). Gate 1. Promoter holding falling three consecutive years with no documented, non-monetising reason is ORANGE, one-notch conviction cap. Gate 2. Promoter holding below 25% in a sub-₹10,000 Cr name with no institutional anchor is ORANGE. Gate 3. Documented open-market promoter buying near CMP is a GREEN corroborating signal — never sufficient alone, never a substitute for Section 5e or 5b-T.

Why / how it binds: Multibaggers in this market are usually promoter-driven compounding machines held for a decade by someone with more at stake than you are. MQF tests the promoter only for the absence of fraud — a floor, not a thesis. This section supplies the missing positive case without touching a single MQF kill-gate or G3 cap.

Section 6 — capex, business model, ROCE decay

Capex-phase names: quantum vs balance sheet; moat-reinforcing vs commodity-capacity test; ROCE recovery timeline; utilisation milestone triggers; capex discipline history; dilution/completion risk. Capex Phase Modifier: board-approved capex, ROCE declining, utilisation <60% → forward ROCE at 75% utilisation, suspend the ROCE-decay signal while documented and moat-reinforcing. Utilisation >70% and ROCE still not recovering → exception expired, activate moat-decay check.

Section 6A Industry Phase & Operating-Leverage Inflection (v2.1 NEW — entry timing)**NAVY**

Industry phase. Classify EMERGENT / INFLECTION / MATURE / DECLINING with evidence. Multibagger Candidate tier requires EMERGENT or INFLECTION, or a documented share-gainer in a MATURE industry (3 consecutive years of revenue CAGR above sector). MATURE + share-flat caps at Steady Compounder. DECLINING caps at WATCH irrespective of valuation — this is the cigar-butt cell and it is not what this suite is for. Operating-leverage inflection. Where the Part 7 N1 bridge relies on a margin delta, name the mechanism and evidence it: capacity utilisation trajectory, gross-block-to-sales turn, fixed-cost absorption, falling interest cost post-deleveraging, or mix shift with segment data. An unevidenced margin delta is set to zero in the bridge and the bridge re-run.

Why / how it binds: The first two years of a multibagger are usually margin, not volume — operating leverage on an existing asset base re-rates earnings faster than any revenue story. The raw material for this already existed scattered across the v2.0 Section 4B bridge, Section 6 capex test and Stage 0 T7; this section is where it is evidenced and pointed at the return calculation.

Section 6B Module B sub-test — Runway & Profit Pool (v2.1 NEW)**NAVY**

State: current market cap; the addressable profit pool (industry revenue × a defensible normalised industry margin); the company's current share of it; and the share implied at the Part 7 N1 target market cap, holding the exit multiple fixed. Gate. If the target requires capturing more than 25% of the addressable profit pool, and the company is not already the leader of a consolidating duopoly, this is RED and the target multiple must be reduced until the implied share is defensible. Size band. Record the market-cap band at entry (micro <₹1,000 Cr / small ₹1,000–10,000 Cr / mid ₹10,000–50,000 Cr / large >₹50,000 Cr). Multibagger Candidate tier above ₹50,000 Cr requires an explicit, logged reason the size objection does not apply.

Why / how it binds: Stops the commonest way an otherwise disciplined return bridge becomes fantasy — a plausible revenue CAGR that, compounded for seven years, implies a market position nobody in that industry has ever held. Size is also the single strongest known conditioner of multibagger base rates, and the pre-v2.1 suite did not condition on it at all.

Section 7 — four-question classifier

(Q1) same moat sources as the core? (Q2) management track record in this area? (Q3) capital deployed above cost of capital? (Q4) domain competence AND succession depth beyond the founding team? Four YES = Accretive Adjacency. Two or more NO = Pivot Risk. Capital below cost of capital for 2+ years = Di-worse-ification. Q4 capability-YES-but-succession-NO = AMBER, linked to 5d I-4.

ROCE trend	Read	Action	Band
–1 to –3pts/1yr	Normal operating variance	Note; watch	TEAL
–3pts/2yrs (not capex)	Early decay signal	Mandatory review; re-run Section 5	ORANGE
–5pts/2yrs	Structural concern	Reduce to 50%; reassess moat	ORANGE
–8pts/2yrs, or ROCE<10%	Thesis likely broken	Exit unless documented capex-phase	RED

ROCE trend	Read	Action	Band
ROCE turns negative	Thesis broken	Immediate exit unless a G2 turnaround was set at entry	RED

Section 9A — the pre-mortem (inside the M5 adversarial protocol)

Mandatory before any verdict is locked, including every rejection. Step 1 steelman the short — strongest bear case as if arguing for payment. Step 2 disconfirming-evidence search, logged, one hostile source minimum. Step 3 pre-mortem paragraph, 4–5 sentences, written before the verdict — it is three years later and this call was wrong; what proved it wrong, which metric moved first. Logged verbatim, re-read first at every forward check. Step 4 48-hour cooling period before first purchase. A report without Steps 1–3 is incomplete regardless of composite; a purchase without Step 4 is a logged process breach.

Modules A/B/C, tiers and sizing

Module A economic moat and quality (now inclusive of 5b-T trajectory and 5f alignment) · Module B growth and scalability (now inclusive of the 6B runway/profit-pool sub-test) · Module C capital efficiency and valuation, cross-checked against the Stage 1 EFV. Each scores to 100. Module C below 65 gates to WATCH regardless of composite.

Condition	Tier	Size	Band
Composite 80+, A 80+, C 80+, all Part 7 gates clear	MULTIBAGGER CANDIDATE — triple conviction	6–8%	GREEN
Composite 76+, A 75+, C 75+, all Part 7 gates clear	MULTIBAGGER CANDIDATE — quality + future	5–6%	GREEN
Composite 70+, A 75+	STEADY COMPOUNDER	4–5%	GREEN
Composite 65–74, C 65+	SPECULATIVE ACCUMULATION	2%	ORANGE
Composite 60–74, C below 65	WATCHLIST	0–1%	ORANGE
Below 60, or G5 review fails, or Stage 1 verdict AVOID/WATCH	SKIP / CAP	0–1.5%	RED
MQF ELEVATED GOVERNANCE RISK or AVOID	GOVERNANCE-CAPPED — binding	0–1%	RED
G3-A	INSTANT SKIP	0%	RED

Multibagger Candidate tier is not available on composite score alone — it additionally requires every Part 7 gate to clear (see Part 7's consolidated gate table). A name that scores 80+ but fails, say, the runway gate or the earnings-dominance test is capped at Steady Compounder. Rule #26 (simplified, v2.1): VIX-based de-sizing is not used and the four-band table is retired (M2 ledger). Tranche for timing only; pause new entries briefly only if liquidity itself is seizing.

Part 4 — Management Quality & Forensic Governance, MQF v1.1

Unchanged from v2.0. Governance is assessed before valuation and can veto Stage 2 sizing outright. A cheap — or fast-compounding — stock with a compromised promoter is not a value pick or a multibagger; it is a slow-motion capital-loss event.

Kill-gates — automatic disqualification regardless of composite

KILL-GATE		
Auditor resignation citing inability to access information or non-cooperation, unresolved		
SEBI debarment of a promoter or director from the securities markets, currently in force		
Promoter or group entity classified as a wilful defaulter, currently in force		
Forensic audit ordered by a regulator or lenders, in progress or with adverse findings		
Conviction — not merely charge-sheet — of a promoter in a fraud-related criminal case		
Verified diversion of company funds to promoter entities, by regulatory or judicial finding		
Pillar	Weight	Source
Related-party transaction conduct	25%	Section 7 (5-yr lookback, minority-disadvantage test)
Board and audit independence	20%	Sections 9, 10
Ownership and group structure transparency	15%	Sections 2, 4, 6
Capital allocation discipline	10%	Section 8
Management stability and remuneration fairness	10%	Sections 11, 12
Leverage, pledge and regulatory-legal cleanliness	10%	Sections 13, 14, 15
Earnings quality (forensic checklist)	10%	Section 17

Each pillar scores 0–10. Composite = $\sum(\text{pillar} \times \text{weight}) \times 10$, out of 100. Section 20A Pattern Escalation: 3+ DISCLOSED/UNRESOLVED or ALLEGATION/MEDIA-CLAIM items clustering within a rolling 18 months and spanning ≥ 2 of Sections 7/9/10/15/16 caps the assessment at ELEVATED GOVERNANCE RISK regardless of composite. Single-pillar floor: any one pillar $\leq 3/10$ caps the verdict at ACCEPTABLE WITH MONITORING even where the weighted composite would reach HIGH-TRUST. Both caps may apply; the lower governs.

Score	Verdict	Portfolio implication	Band
85–100	HIGH-TRUST MANAGEMENT	No governance-based size constraint	GREEN
65–84	ACCEPTABLE WITH MONITORING	Position size capped; re-run RPT/pledge/litigation every 2Q	ORANGE
45–64	ELEVATED GOVERNANCE RISK	Stage 2 sizing capped at WATCHLIST. Tracking position at most	RED
Below 45, or any kill-gate	AVOID	Not investable irrespective of valuation, growth or engine score	RED

Section 19 tags: VERIFIED FACT (Tier 1–2, cite document and date) / DISCLOSED-UNRESOLVED (state current stage) / ALLEGATION-MEDIA CLAIM (Tier 5, name source, never present with the confidence of a verified fact).

Part 5 — Portfolio & Execution Layer, v2.1

Nothing here overrides a governance cap; everything here can reduce a position below what Stage 2 would otherwise allow. v2.1 rewrites P-1, P-5, P-6 and P-7 — the rules that previously capped a working compounder on its way up.

Rule	Content	Band
P-1 (v2.1 amended)	10–15 positions. Max 8% at cost — unchanged, never relaxed. Cap at market raised to 25% for a position whose moat trajectory (5b-T) is BUILDING or HOLDING and whose ROIIC (5e) remains above cost of capital. Above 25% at market: trim review, not automatic trim. Adding above the 8% cost cap remains prohibited in all cases. Minimum meaningful position 2%.	NAVY
P-2	Max 25% of equity capital in any one sector; max 35% in any one macro theme, declared at entry.	TEAL
P-3	Correlation clustering: name the dominant macro factor before entry; if it already exceeds 30% of the book, halve or defer the new position.	TEAL
P-4	Cash is a residual, not a market call. Cash below 5% while 3+ names sit in the entry queue signals sizing is too aggressive.	TEAL
P-5 (v2.1 re-ranked)	Queue rank order amended to: Part 7 N1 bridge expected return net of E-3 → Section 5e engine growth → Stage 2 composite → MoS margin. Unfunded names logged DEFERRED with a re-check date, never quietly dropped.	NAVY
P-6 (v2.1 two-sided)	Position down >30% from cost: mandatory thesis re-read (pre-mortem first, per M5), written hold/add/exit. New in v2.1: position up >100% also triggers a mandatory written re-read — hold/trim/add — because selling a working thesis for the wrong reason is the commonest multibagger failure mode and deserves the same documentation as a loss.	NAVY
P-7 (v2.1 replaced — closes diagnosis D-6)	v2.0 text retired for engine-qualified names: "trim 20–25% at base-case fair value with the thesis intact" forced an exit exactly where a multibagger's return is made. Compounder Hold Rule (governs for engine-qualified names): trim or exit only on (a) thesis-break per the journal entry; (b) ROIIC below cost of capital for two consecutive years (5e Gate 2); (c) moat trajectory turning ERODING (5b-T); (d) price exceeding 2× the N1 exit-multiple assumption on delivered earnings; (e) a P-2/P-3 concentration limit binding. Base-case fair value is a reference, not a trigger. The v2.0 P-7 ladder stands unchanged for non-engine names.	NAVY

Execution rules E-1 to E-4 (unchanged)

Rule	Content
E-1 Entry	Rule #26 tranching for timing only. Limit orders in names below ₹5 Cr ADTV. No entry on result/corporate-action day.
E-2 Liquidity floors	Days-to-exit = position value ÷ (25% of 30-day ADTV), must be ≤5 trading days. Cut size until it is, or do not enter.
E-3 Cost drag	Estimate round-trip cost (brokerage + STT + realistic slippage for the ADTV band) and subtract from the M7 expected excess return before finalising the verdict.
E-4 Tax	State returns net of LTCG/STCG by expected holding period.

Review cadence: Quarterly — P-2/P-3 exposure, DEFERRED queue, any P-6 30%-down write-up. Semi-annual — MQF Sections 7/13/14/15/16 on every holding. Annual — Rule #27 scorecard, Section 5d re-score, M1 denominators, M2 audit, M4 kill-switch test, M8 funnel hit-rate.

Part 7 — The Multibagger Engine (v2.1 NEW)

Part 7 is the orchestration layer. Sections 5b-T, 5e, 5f, 6A and 6B (Part 3) supply the evidence; Part 7 turns that evidence into three cross-cutting tests — the return bridge, the variant-perception claim, and the funnel screen that feeds all of it — and then a single consolidated gate. A name that fails the gate cannot hold Multibagger Candidate tier at any composite score. A name that passes is not thereby a BUY: every Stage 1 trap test, every Irani flag, the MQF verdict and every Part 5 limit still apply in full. Placed here, ahead of Part 6, because Part 6 packages this Part's output into the final report.

N1 The Multibagger Return Bridge (mandatory, binding)	NAVY
State the target multiple and horizon, convert to a required annualised return, decompose into four named, separately evidenced terms: (a) revenue CAGR (anchored to capacity, order book or industry volume); (b) margin delta (mechanism named per Section 6A); (c) share-count change (fully diluted — warrants, ESOP pool, announced QIP); (d) exit multiple — the re-rating term. Total return $\approx (1 + \text{EPS CAGR}) \times (1 + \text{re-rating CAGR}) - 1$. N1-1 Earnings dominance. At least 60% of the required log-return must come from earnings, not re-rating; a bridge relying mostly on the multiple is a re-rating trade and cannot hold Multibagger Candidate tier. N1-2 Exit-multiple ceiling. The assumed exit multiple may not exceed the lower of the name's own and the sector's 10-year 75th-percentile multiple; no exit multiple above entry where Section 5b-T trajectory is ERODING. N1-3 Dilution honesty. Computed on fully diluted share count; a bridge closing only undiluted is ORANGE and caps at ACCUMULATE. Why / how it binds: This is the arithmetic the pre-v2.1 suite never wrote down. Most names that feel like multibaggers cannot produce a bridge that closes without an unsupportable growth or exit-multiple term — the bridge makes that visible in minutes rather than after hours of scoring, and gives the Decision Journal a falsifiable, dated claim to grade.	

N1 reference — required CAGR by target and horizon

Target	3 years	5 years	7 years	10 years
3x	44.2%	24.6%	17.0%	11.6%
5x	71.0%	38.0%	25.8%	17.5%
10x	115.4%	58.5%	39.0%	25.9%

Worked example — 5x in 5 years (38.0% required). Entry P/E 15, exit P/E 25 $\rightarrow$ re-rating = $(25/15)^{(1/5)} - 1 = 10.8\%$ p.a. Required EPS CAGR = $1.380/1.108 - 1 = 24.6\%$ p.a. Earnings share of log-return = $\ln(1.246)/\ln(1.380) = 68\% \rightarrow$ clears N1-1. At 4% p.a. dilution, required net-profit CAGR rises to $\sim 29.6\%$ — the figure Section 6A's revenue CAGR $\times$ margin delta must actually deliver.

N7 Variant Perception & Implied Expectations (closes diagnosis D-10)	NAVY
Answer four questions in writing, before the verdict: (1) what does the market currently believe? — evidenced with a reverse DCF, solving for the growth and duration the current price embeds (distinct from Section 4C's forward DCF); (2) what specifically is wrong with that belief, with Tier 1-3 evidence; (3) why does the mispricing persist? — name the mechanism, which must be one of the three M3 edges: coverage neglect / time-horizon mismatch / structural non-participation (size, liquidity, mandate) / optical ugliness (trough earnings, capex drag, ex-parent complexity); (4) what closes the gap, and by when — a dated, observable event or metric, not "rerating". A name that cannot answer all four is capped at WATCH, per M3. Why / how it binds: Replaces the "catalyst score" the v2.0 output format referenced but never defined. Not how many catalysts exist, but what the market has wrong and why it has stayed wrong — giving the Section 9A pre-mortem something real to attack and the Decision Journal something falsifiable to grade.	

N8 The Multibagger Funnel Screen (closes M8's unspecified systematic screen)**NAVY**

Run monthly on the full listed universe, logged with date. Filters: market cap ₹300 Cr-₹15,000 Cr; revenue CAGR (3-yr) $\geq 15\%$ or a documented inflection from a loss/low base; ROCE $\geq 15\%$ and rising, or crossing $10\% \rightarrow 15\%$ for the first time; ROIC (3-yr) $\geq 18\%$ (the Section 5e pre-filter); 3-yr average payout $\leq 40\%$; D/E < 1.0 and interest coverage $> 3.0\times$; promoter holding $\geq 40\%$ and not falling; institutional holding $< 15\%$ (screening deliberately for coverage neglect); OCF/EBITDA 3-yr average > 0.60 ; 30-day ADTV $\geq ₹1$ Cr (G0). Ranked by $ROIC \times (1 - \text{payout})$; top 20 enter the funnel with candidate source recorded per M8. This is a funnel, never a verdict — every name still runs the full suite.

Why / how it binds: M8 identified the funnel as the dominant driver of outcomes and left the input undefined, so it stayed news-driven — which is how a value-discipline framework ends up looking at value-shaped names. This screen is explicitly biased toward reinvestment, neglect and inflection, because that is where the base rate lives, and is directly implementable in the existing Nifty 500 Colab screener.

The Multibagger Candidate Gate — consolidated (governs Stage 2 Modules table and Part 6 matrix)

Multibagger Candidate tier requires every gate below to clear. A single failure caps the name at Steady Compounder (one soft failure) or Speculative Accumulation/WATCH (a hard failure, RED row) — never a route back to full tier by strength elsewhere.

Gate	Source	Pass condition	Band on failure
Engine growth clears requirement	Section 5e	$g(ROIC \times RR) \geq N1$ required net-profit CAGR	RED — Steady Compounder cap
ROIC floor	Section 5e	ROIC above cost of capital, 2 consecutive years	RED — hard fail, Speculative/WATCH
Payout ceiling	Section 5e	3-yr average payout $\leq 40\%$	RED — hard fail, routes to Steady Compounder
Moat trajectory	Section 5b-T	BUILDING on ≥ 1 moat source, Tier 1-3 evidence	ORANGE — Steady Compounder cap
Runway / profit pool	Section 6B	Implied share at target $\leq 25\%$ of addressable pool, or documented consolidating-duopoly leader	RED — hard fail, reduce target
Industry phase	Section 6A	EMERGENT/INFLECTION, or documented share-gainer in MATURE	RED (DECLINING) — WATCH regardless of valuation
Promoter alignment	Section 5f	No open Gate 1/2 flag, or flag explained and logged	ORANGE — one-notch conviction cap
Return bridge closes	Part 7 N1	Earnings share of log-return $\geq 60\%$; fully diluted; exit multiple within N1-2 ceiling	RED — hard fail, re-run at reduced target
Variant perception complete	Part 7 N7	All four questions answered, reverse DCF run, mechanism is a named M3 edge	RED — hard fail, caps at WATCH per M3
MQF governance	Part 4	HIGH-TRUST (85+) for the Engine Compounder MoS tier; ACCEPTABLE (65+) minimum for any Multibagger sizing	RED — binding governance cap, never outvoted

M1 status: every gate above is new to v2.1 and enters at PILOT. No report may describe a gate as VALIDATED or PROVISIONAL. Conviction is reduced one notch per Standing Rule 3 for any BUY carrying a PILOT gate — in year one, that is every Multibagger Candidate BUY: none may be logged at HIGH conviction until denominators exist (see the M2 retirement ledger and validation protocol at the end of this document).

Part 6 — Output Format, Checklist and Integration

The integration matrix, the report structure, and the checklist all change to accommodate Part 7. Everything else in this Part carries forward from v2.0 unchanged.

The integration matrix — Engine × Price (v2.1 replaces the Cheap × Durable 2×2, closes D-1)

The primary axis is now the compounding engine (the Part 7 gate); the secondary axis is price versus the return the engine can deliver — the N1 bridge, net of E-3 costs and the M7 hurdle. "Cheap" as a standalone concept no longer gates anything on its own.

	Price BELOW bridge (room to spare)	Price AT bridge (clears, no cushion)	Price ABOVE bridge (does not clear)
ENGINE STRONG all gates pass	BUY — Multibagger Candidate tier, size per matrix and Cap Reconciliation	BUY at one tier lower — accumulate on weakness, staged per Rule #26. This is the cell v2.0 sent to WATCH, and it is where most realised multibaggers were buyable.	WATCH — named, dated price condition; re-check quarterly
ENGINE PARTIAL one soft gate unmet	ACCUMULATE — capped at Steady Compounder until the gate is met	WATCH — named, dated evidence condition on the failing gate	WATCH / IGNORE per Stage 1 composite
ENGINE WEAK any hard gate fails	Value candidate only — runs the Stage 1 value path, never Multibagger tier, never above 3%	AVOID	IGNORE

Conservative-governs applies only to governance RED flags and the MQF cap, unchanged from v2.0. PARTIALLY on either axis still defaults to WATCH; an upgrade requires a named, dated condition that would move that axis to YES.

Report structure — mandatory order

#	Block	Must contain
0	Colour legend	Reproduced in full; every finding assigned exactly one band
1	Stock snapshot	Name, ticker, sector, CMP, mcap, analysis date; macro anchors dated and cited; candidate source per M8/N8
2	Part 1 — Stage 0	Triggers checked; cluster validation composition (S11); tracking-position status if S1 invoked; X1-X4; re-screen condition and date
3	Part 2 — Stage 1	Gates, Irani flags, Layer A grid with MATM anchor and trajectory note, Rule #25 line, Layer B/EFV/bear-case gate, composite, M7 hurdle, verdict, attractive-below
4	Part 3 — Stage 2	Section 5 moat + 5b-T trajectory, 5c decay, 5d table, 5e engine growth vs N1 requirement, 5f alignment, 6/6A/6B, Section 7, Modules A/B/C, M5 protocol output printed before the verdict, G3 state
4A	Part 7 — Multibagger Engine	N1 return bridge with all four terms and earnings-share calc; N7 four variant-perception answers incl. reverse DCF; N8 candidate source; the consolidated gate table with each row's pass/fail
5	Part 4 — MQF (when run)	Kill-gate status; pillar scores; Section 20A status; raw AND effective capped verdict
6	Part 5 — Portfolio	P-2/P-3 exposure impact; P-5 queue rank; M3 edge relied upon; Compounder Hold Rule conditions if engine-qualified (block 6A)
7	Combined final recommendation	The Engine×Price matrix placement and the final call block
8	Sources	Every live figure with source, date, M6 tier
9	Condensed summary, then Stage 0 findings	In that order, by default, without being asked

Checklist Q58–Q88

#	Check
Q58-Q64	G-Sec/VIX dated; moat-decay scorecard done; reverse DCF (N7) + probability-weighted EFV (4C) both cited and not conflated; catalyst→variant-perception (N7) done; days-to-exit ≤5 (E-2); journal entry with review dates
Q65-Q70	Section 5 moat classified; Section 6/capex done; Section 7 answered; compounder tier stated; Section 5d rated; pre-mortem logged before verdict
Q71	Section 1A scope correctly applied — full Stage 2 only where v2.1's amended rule requires it, not blanket (N/A = qualitative-only rejection)
Q72-Q73	MQF checked, cap applied if required; Section 5d cross-checked against Rule #27/Section 6, downgrade applied if triggered
Q74-Q79	M1 status tags correct; M3 edge named; M5 steelman/search/cooling done; M6 tiers stated, no Tier-4-only GREEN; M7+E-3 stated; M8/M9 candidate source and both composites logged
Q80	Version sync: instruction set names the same component versions as this PDF (v2.1: Stage 0 v1.3 · Stage 1 v3.5 · Stage 2 v7.0 · MQF v1.1)
Q81	N1 bridge built, four terms evidenced, earnings share ≥60% stated, fully diluted share count used?
Q82	N1-2 exit-multiple ceiling checked against both own and sector 10-yr 75th percentile?
Q83	Section 5e ROIC, RR and engine growth computed and compared against the N1 requirement? Payout ≤40% for Multibagger tier?
Q84	Section 6B profit pool stated, implied share at target computed, gate tested, size band recorded?
Q85	Section 5b-T trajectory rated per source at Tier 1-3; MATM conditioning applied correctly?
Q86	Section 6A industry phase classified; margin-delta mechanism named and evidenced, or zeroed and bridge re-run?
Q87	N7 all four questions answered; persistence mechanism mapped to a named M3 edge?
Q88	N8 candidate source recorded; is the 50% systematic-screen share (M8) still being met year-to-date?

Final call block

FINAL CALL: BUY / ACCUMULATE / WATCH / AVOID
 Conviction: HIGH / MEDIUM / LOW (−1 notch per PILOT rule invoked; all Part 7 gates are PILOT)
 Engine verdict: STRONG / PARTIAL / WEAK — consolidated gates passed [x of 9]
 Return bridge (N1): [target]x / [n]yrs = [r]% p.a. → EPS CAGR [e]% × re-rating [m]% | earnings share [p]%
 Engine growth (5e): ROIC [i]% × RR [rr]% = [g]% vs [e]% required — CLEARS / FAILS
 Moat trajectory (5b-T): BUILDING / HOLDING / ERODING → MATM anchor applied [1.00/1.10/1.25/1.50]x
 Compounder tier: Multibagger Candidate / Steady Compounder / Speculative Accumulation / Watch
 Position: Stage-1 cap [X]% | Stage-2 base [Y]% | MQF cap [G]% | Part 5 limit [P]% | Reconciled [Z]%
 Edge relied on (M3): time-horizon / coverage-neglect / behavioural
 Why mispriced (N7): [persistence mechanism, one line]
 Benchmark hurdle (M7): index, required return, driver — net of E-3
 Attractive below: Rs [X]
 Hold rule: Compounder Hold Rule (a)-(e) with the trigger number for each, or v2.0 P-7 ladder if non-engine
 Thesis-breaker/exit: [one line]
 PILOT rules invoked: [list]

Revised workflow order

STEP 0 Stage 0 (incl. S1 tracking carve-out) → STEP 0A Part 7 N1/N7/N8 pre-screen: build the bridge, run the gate table, answer variant perception — a name whose bridge cannot close, or that fails a hard gate, stops here with a note and a pre-mortem, which is the release's largest time saving → STEP 1 read this PDF → STEP 2 Stage 1 Pass 1 at MATM 1.00x → STEP 3 Stage 2 per the amended Section 1A scope, including 5b-T/5e/5f/6A/6B → STEP 3A Stage 1 Pass 2 with the trajectory-conditioned MATM → STEP 3B MQF → STEP 3C Part 5, including the Compounder Hold Rule where engine-qualified → STEP 4 deliver in the Part 6 order above.

M2 Retirement Ledger — how v2.1's new rules were paid for

Retired / amended	Reason	Replacement
Section 1A blanket full-coverage rule	Conflicted directly with M8: spent the binding constraint (analyst hours) on rejected names.	Scoped rule (Part 3): full Stage 2 on FAIRLY VALUED+ and price-only rejections; qualitative note + pre-mortem elsewhere.
Rule #26 four-band VIX table	The rule's own text already said VIX is not used for de-sizing — the table was decorative and, per M2's own audit standard, never changed a verdict.	One line: tranche for timing only; pause new entries only if liquidity itself is seizing.
Layer A dividend-yield metric (engine-assessed names only)	Rewards the exact behaviour Section 5e Gate 3 identifies as incompatible with compounding; already double-counted in the Total-Yield Tolerance Band.	Retired for Part 7 names; retained in full for the non-engine value path.
Total-Yield Tolerance Band — nominated, not yet executed	PILOT at n=1, circular (ITC as both calibration and live case).	Carried forward unchanged this release; formally nominated for deletion at the next annual audit unless a third independent case is logged first.

The debt that remains outstanding. Three retirements do not fully pay for the ten new/amended rules this release adds (Section 5b-T, 5e, 5f, 6A, 6B, N1, N7, N8, plus the P-1/P-5/P-6/P-7 rewrites). M2's other route is a denominator, and none of the new rules has one. Every Part 7 gate and every new Stage 2 section therefore enters at PILOT under M1, stated plainly rather than smoothed over.

Validation protocol — denominators to collect before any N-rule can leave PILOT

Rule	Denominator required	First review
N1	N bridges built; N that closed; of those, N delivering $\geq 50\%$ of bridged EPS CAGR at 24mo	24 months
Section 5e	N names passing the gate; realised 3-yr ROIIC vs assumed at entry; N short by $>30\%$	24 months
Section 6B	N names where the gate fired; realised market share at 36mo vs implied share	36 months
Section 5b-T	N names BUILDING at entry; N still BUILDING at 24mo; return spread by trajectory cohort	24 months
N8	N screened/month; N entering the funnel; N reaching BUY; hit-rate vs the news-sourced funnel	12 months
S1/S12	N tracking positions opened; N reaching evidence date intact; N exited at a loss first	24 months, auto-revert on failure

Standing honesty note

Under M1, no rule in this suite — old or new — currently holds VALIDATED status. The two levers in this release that matter most are not rules: they are the funnel (N8) and the holding period (the Compounder Hold Rule at P-7). Running the N8 screen monthly and holding through the Compounder Hold Rule captures most of the available improvement even without engaging the rest of Part 7 in depth; implementing all of Part 7 while still trimming winners at base-case fair value captures none of it.

What v2.1 does not fix. It supplies no denominators — every N-rule must be treated as PILOT in visible output. It also does not resolve a deeper tension: a framework this conservative on governance will systematically exclude some genuine multibaggers, because promoter-driven smallcaps mid-compounding are frequently imperfect on exactly the axes MQF measures. That exclusion is deliberate and correct — the tail risk avoided is permanent capital loss — but it is a real source of false negatives and should be measured, not assumed away: add to the annual audit what happened to names excluded by Section 20A pattern escalation alone, with no kill-gate fired.

The failure mode to watch in v2.1. v2.0's failure mode was buying nothing. v2.1's risk is the opposite, disguised as rigour: the N1 bridge is a construction, and a motivated analyst can make it close by nudging the revenue CAGR two points and the exit multiple three turns. N1-1's earnings-dominance rule, N1-2's

percentile ceiling, and the Section 6B profit-pool gate make that harder, not impossible. The discipline that protects the bridge is M5's pre-registration habit: write the bridge inputs into the Pre-Commitment Log with a timestamp before computing the output value, exactly as Section 4C already requires for probabilities. A bridge computed after the desired answer is known is not evidence — it is arithmetic-shaped advocacy.

And the thing no framework can supply. A multibagger is mostly a holding-period phenomenon. The N1 arithmetic assumes you are still holding in year five, through at least one 40% drawdown every one of these names will have. No rule anywhere in this suite makes that easier. The M3 time-horizon edge is the primary claimed edge for a reason, and it is the only one of the three entirely within your control.

Revision history

Version	Date	Summary
Suite v2.0	Sep 2026	Red-team hardening release. Part 0 Meta-Layer (M1-M9), Part 5 Portfolio & Execution Layer, 24-finding register. Trap-avoidance architecture completed.
Suite v2.1	Sep 2026	MULTIBAGGER CAPABILITY RELEASE — supersedes v2.0 in full. Diagnosis (D-1 to D-11) found the integration order, the missing return arithmetic, the absent reinvestment-rate test, the static moat test and the exit rules jointly prevented the suite from holding a multibagger. NEW Part 7 — N1 return bridge, N7 variant perception + reverse DCF, N8 funnel screen, consolidated candidate gate — orchestrating five new Stage 2 sections (5b-T moat trajectory, 5e compounding engine, 5f promoter alignment, 6A industry phase, 6B runway/profit pool). Section 1A scoped to stop conflicting with M8. Integration matrix inverted to Engine×Price. P-1 raised to 25% at market for engine-qualified names; P-7 replaced by the Compounder Hold Rule; P-6 made two-sided; P-5 re-ranked. Stage 1 adds the MQF-gated Engine Compounder MoS tier. Stage 0 adds the S1/S12 tracking carve-out with an auto-reverting sunset. M2 paid by retiring/simplifying Section 1A, Rule #26's VIX table and the Layer A dividend-yield metric. Checklist extended to Q88. All new rules enter at PILOT under M1.

For personal investment research and education ONLY. Not investment advice, not a SEBI research report, not a recommendation. All equity investment carries market risk including full loss of principal. The framework organises research; it does not guarantee any outcome. Verify all live data independently and consult a SEBI-registered investment adviser before acting. Branding: STOCK RESEARCH only.